


GitOps Failure Modes: When ArgoCD Sync Breaks



Published on June 5, 2022



Webstack
Builders

Table of Contents

Understanding ArgoCD Sync Mechanics	4
The Sync Process Explained	4
Sync Waves and Resource Ordering	5
Sync Hooks Explained	7
Common Sync Failure Categories	9
Resource Dependency Failures	9
Hook Failures	11
Health Check Failures	13
Diff/Drift Detection Issues	15
Debugging Workflow	17
Step 1: Identify the Failure Point	17
Step 2: Examine ArgoCD Logs	18
Step 3: Inspect Kubernetes State	19
Step 4: Reproduce and Fix	20
Specific Failure Scenarios	21
Scenario: Sync Stuck on Hook	21
Scenario: OutOfSync Loop	23
Scenario: CRD Sync Order Problem	25
Scenario: Resource Pruning Gone Wrong	26
Prevention Strategies	28
Manifest Validation	28
Sync Windows	30
Monitoring and Alerting	31
Advanced Troubleshooting	33
Resource Tracking Methods	33
Debugging Manifest Generation	34
Recovering from Corrupt State	35
Conclusion	36



GitOps (Git as single source of truth for declarative infrastructure) sells you on declarative simplicity: define your desired state in Git, and ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) makes it so. The pitch works beautifully in demos. In production, you'll eventually stare at a sync that's been "Progressing" for 20 minutes, wondering what's actually happening inside the black box.

I hit this wall during a production deployment that had worked flawlessly in staging for months. The sync started normally, applied the first few resources, then stopped. No error. No timeout. Just... stuck. The ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) UI showed "Syncing" with a spinner. Kubernetes showed pods running. Logs showed nothing useful. Manual `kubectl apply` worked fine – it only hung when deployed through ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes).

The root cause turned out to be a sync wave ordering issue. A PreSync hook was waiting for a service that wouldn't exist until wave 0, but the hook was in wave -1. In staging, the service happened to exist from a previous deployment. In production, we'd just created the namespace. The dependency was always broken – we just never noticed because we never deployed to a clean environment.

This is the gap between GitOps (Git as single source of truth for declarative infrastructure) theory and operational reality. The declarative model abstracts away the *how* of deployment, which is great until something goes wrong. Then you need to understand exactly what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) is doing: how it orders resources, when hooks execute, how it determines health, and why it might decide your perfectly valid manifests can't be applied.

WARNING

GitOps (Git as single source of truth for declarative infrastructure) hides complexity until sync fails. When ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) gets stuck, you need to understand what's happening underneath the declarative abstraction to debug it.

This article breaks down the sync process, catalogs the failure modes I've encountered (and how to recognize them), and provides a systematic debugging workflow. The goal isn't to memorize every edge case – it's to build a mental model of what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) is actually doing so you can reason about failures when they happen.



Understanding ArgoCD Sync Mechanics

Before debugging failures, you need a clear picture of what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) is actually doing during a sync. The process has more steps than most people realize, and failures can occur at any of them.

The Sync Process Explained

When you push a commit or click “Sync” in the UI, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) kicks off a multi-phase process. Understanding these phases tells you where to look when something goes wrong.


First, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) detects the change – either through polling the Git repository or receiving a webhook. It then generates the manifests, which might involve running Helm template, Kustomize build, or just reading raw YAML (YAML Ain't Markup Language) files. This is where Helm value resolution and Kustomize overlays get applied.

Next comes diff calculation. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) compares the generated manifests against the live cluster state. If there's no difference, you're already synced – nothing to do. If there is a difference, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) needs to figure out *how* to apply the changes.

Resource ordering happens next. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) sorts resources by sync wave, then by kind (namespaces before deployments, for example), then by name. This ordering determines which resources get applied first.

Finally, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) applies resources wave by wave, runs any hooks at the appropriate phases, and waits for each resource to become healthy before moving on. If any resource fails to become healthy within the timeout, the sync fails or marks the application as degraded.





ArgoCD sync process from Git commit to healthy application state

Sync Waves and Resource Ordering

Kubernetes doesn't guarantee the order in which resources are created. If you `kubectl apply -f a` directory, resources might be created in any order. Usually that's fine – Kubernetes' reconciliation loops eventually sort things out. But "eventually" can mean failed pods, restarts, and transient errors that confuse monitoring.

Sync waves give you explicit control over ordering. Each resource can have an `argocd.argoproj.io/sync-wave` annotation with an integer value. Lower numbers go first. Resources without the annotation default to wave 0.



manifests/app/resources.yaml

```

1  # Kubernetes manifests with ArgoCD sync-wave annotations.
2  # These annotations are added directly to your application manifests
3  # in your GitOps repository. ArgoCD reads them during sync to determine
4  # resource application order.
5
6  # Wave -1: Secrets and ConfigMaps first
7  apiVersion: v1
8  kind: Secret
9  metadata:
10   annotations:
11     argocd.argoproj.io/sync-wave: "-1"
12   name: app-secrets
13   ---
14  # Wave 0: Services and deployments (default)
15  apiVersion: apps/v1
16  kind: Deployment
17  metadata:
18   annotations:
19     argocd.argoproj.io/sync-wave: "0"
20   name: app
21   ---
22  # Wave 1: Post-deployment resources
23  apiVersion: batch/v1
24  kind: Job
25  metadata:
26   annotations:
27     argocd.argoproj.io/sync-wave: "1"
28   name: db-migration

```

Sync wave annotations controlling resource application order.

The wave number itself is arbitrary – what matters is the relative ordering. I’ve seen teams use -10 to 10, and others use -2 to 2. Pick a convention and stick with it. Here’s a pattern that works for most applications:

Wave	Typical Resources	Purpose
-2	Namespaces, CRDs	Infrastructure prerequisites



Wave	Typical Resources	Purpose
-1	Secrets, ConfigMaps	Configuration dependencies
0	Deployments, Services	Main application (default)
1	Jobs, migrations	Post-deployment tasks
2	HPA, PodDisruptionBudgets	Scaling configuration

Common sync wave ordering pattern for application deployments.

INFO

Resources in the same sync wave are applied in parallel. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) waits for all resources in a wave to be healthy before proceeding to the next wave.


The “wait for healthy” behavior is critical. If wave -1 has a Secret and a ConfigMap, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) applies both simultaneously and waits for both to exist before moving to wave 0. If your Deployment in wave 0 references that Secret, you’re guaranteed it exists. Without sync waves, you’d have a race condition.

Sync Hooks Explained

Sync waves control ordering *within* the normal apply process. Hooks let you run resources *outside* that process entirely – before sync starts, after it completes, or when it fails.

The most common use case is database migrations. You want to run migrations after the new code is deployed but before traffic shifts to it. Or you want to take a backup before any changes are applied. Hooks give you those insertion points.





Hook execution sequence during ArgoCD sync lifecycle

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) supports five hook types, each running at a specific point in the sync lifecycle:

#	Hook Type	When It Runs	Use Case
1	PreSync	Before any manifests applied	Database backups, pre-flight checks
2	Sync	During normal sync	Rarely used (default behavior)
3	PostSync	After all manifests healthy	Notifications, smoke tests
4	SyncFail	When sync fails	Cleanup, alerting



#	Hook Type	When It Runs	Use Case
5	Skip	Never (resource skipped)	Templates, documentation

ArgoCD hook types with execution timing and use cases.

Hooks are typically Kubernetes Jobs that run a script or command. You define them as regular manifests in your GitOps (Git as single source of truth for declarative infrastructure) repo, but the hook annotations tell ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) to treat them specially:

manifests/app/migrations/db-migration-job.yaml

```

1  # PreSync hook for database migrations. This Job runs before ArgoCD
2  # applies any other manifests, ensuring the database schema is ready
3  # before the new application version starts. Place in your GitOps repo
4  # alongside your application manifests.
5
6  apiVersion: batch/v1
7  kind: Job
8  metadata:
9    name: db-migration
10   annotations:
11     argocd.argoproj.io/hook: PreSync
12     argocd.argoproj.io/hook-delete-policy: HookSucceeded
13     argocd.argoproj.io/sync-wave: "-1"
14  spec:
15    template:
16      spec:
17        containers:
18          - name: migrate
19            image: app:latest
20            command: ["/migrate.sh"]
21            restartPolicy: Never
22        backoffLimit: 3

```

PreSync hook for database migration with cleanup policy.



The `hook-delete-policy` annotation controls what happens to the hook resource after it runs.

- `HookSucceeded` deletes the Job after it completes successfully, keeping your namespace clean.
- `BeforeHookCreation` deletes the old hook before creating a new one — useful when you want to see the previous hook's logs until the next sync. Without a delete policy, hook resources accumulate, and old hooks can block new syncs if they're still present.

Hooks and sync waves can be combined. In the example above, the migration hook has both `hook: PreSync` and `sync-wave: "-1"`. This means it runs before the normal sync *and* before any other PreSync hooks in higher waves. If you have multiple PreSync hooks that depend on each other, use sync waves to order them.

Common Sync Failure Categories


Now that you understand the mechanics, let's look at the failure modes. I've grouped these into four categories based on where in the sync process they occur. Recognizing which category you're dealing with tells you where to start debugging.

Resource Dependency Failures

The most common sync failures happen because a resource references something that doesn't exist yet. Your Deployment needs a Secret, but the Secret hasn't been created. Your Pod uses a ServiceAccount, but the ServiceAccount's RoleBinding is missing.

These failures are frustrating because they're often intermittent. Sync works sometimes (when resources happen to be created in the right order) and fails other times (when they're not). The randomness makes them hard to reproduce and diagnose.





Resource dependency graph showing potential failure points

When a dependency is missing, pods fail to start with characteristic error messages:

Common Symptoms:

- Pod stuck in `CreateContainerConfigError`
- `secret "x" not found` errors
- `serviceaccount "x" not found` errors

The fix is usually straightforward: put the dependency in an earlier sync wave. But first you need to identify which dependency is missing:

debug-dependency-failure.sh

```
1  #!/bin/bash
2
3  # Check what secrets a deployment needs
4  kubectl get deployment app -o
   jsonpath='{.spec.template.spec.containers[*].envFrom[*].secretRef.name}'
5
```



```

6  # Check if secret exists
7  kubectl get secret app-secrets
8
9  # Check ArgoCD resource status
10 argocd app get myapp --show-operation
11
12 # Check events for failed pods
13 kubectl get events --field-selector reason=Failed --sort-by='.lastTimestamp'

```

Commands for debugging resource dependency failures.

Once you’ve found the missing dependency, the fix is to assign it a lower sync wave than the resource that needs it. But be aware that sync waves only help across different waves – they don’t help within a single wave.

⚠ DANGER

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)’s default sync does not guarantee ordering between resources in the same wave. If resource A depends on resource B and both are in wave 0, the sync may fail randomly depending on apply order.

Hook Failures

Hooks fail differently than normal resources because they’re Jobs, not long-running workloads. A Deployment that can’t start will keep retrying – you have time to notice and fix it. A hook Job that fails blocks the entire sync immediately.

The trickiest hook failures are timeouts. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) waits for hooks to complete before proceeding, and “complete” means the Job reached a terminal state (Succeeded or Failed). If your hook runs longer than expected, the sync just hangs. There’s no timeout by default – ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) will wait forever.

Failure Mode	Symptom	Cause
Hook timeout	Sync stuck at "Running PreSync hooks"	Hook job takes too long



Failure Mode	Symptom	Cause
Hook crash	Sync failed, hook pod in CrashLoopBackOff	Bug in hook script
Hook not deleted	Old hook blocks new sync	Missing delete policy
Hook wrong phase	Resources applied before prerequisites	Wrong hook type

Common hook failure modes with symptoms and causes.

Table: Common hook failure modes with symptoms and causes.

For long-running hooks like database migrations, always set `activeDeadlineSeconds` on the Job. This gives ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) a timeout – if the migration takes longer than expected, the sync fails instead of hanging forever:

manifests/app/migrations/slow-migration-job.yaml

```

1  # Long-running migration hook with explicit timeout.
2  # The activeDeadlineSeconds ensures ArgoCD doesn't wait forever
3  # if the migration hangs or takes longer than expected.
4
5  apiVersion: batch/v1
6  kind: Job
7  metadata:
8    name: slow-migration
9    annotations:
10     argocd.argoproj.io/hook: PreSync
11     argocd.argoproj.io/sync-wave: "-1"
12  spec:
13    activeDeadlineSeconds: 600 # 10 minute timeout
14    template:
15      spec:
16        containers:
17          - name: migrate
18            image: app:latest
19            command: ["/long-migration.sh"]
20        restartPolicy: Never

```



Hook configuration with extended timeout for long-running operations.


The “hook not deleted” failure deserves special mention. If your previous sync left a hook Job lying around (because you didn’t set a delete policy), and that Job has the same name as the new hook, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) can’t create the new one. The fix is to add `hook - delete-policy: BeforeHookCreation` so ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) cleans up old hooks before creating new ones.

Health Check Failures

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) doesn’t just apply resources – it waits for them to become healthy. “Healthy” means something different for each resource type. Deployments are healthy when their replicas are ready. Services are healthy when they have endpoints. Jobs are healthy when they complete successfully.

For built-in <https://kubernetes.io/docs/concepts/workloads/> Kubernetes types, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) has hardcoded health checks that usually do the right thing. The problems start when you use Custom Resources. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) doesn’t know how to assess the health of your custom CRD (Custom Resource Definition), so it just marks it “Healthy” immediately after creation – even if the controller hasn’t reconciled it yet.





ArgoCD health assessment flow for different resource types

For CRDs, you need to write custom health checks in Lua.¹ These go in the `argocd-cm` ConfigMap

argocd-cm-health-customization.lua

```

1  -- Custom health check for a CRD.
2  -- Add to argocd-cm ConfigMap under:
3  --   resource.customizations.health.<group>_<kind>
4  -- Example key: resource.customizations.health.example.com_MyResource
5
6  hs = {}
7  hs.status = "Progressing"
8  hs.message = ""
9
10 if obj.status ~= nil then

```



```
11     if obj.status.phase == "Ready" then
12         hs.status = "Healthy"
13         hs.message = "Resource is ready"
14     elseif obj.status.phase == "Failed" then
15         hs.status = "Degraded"
16         hs.message = obj.status.message or "Resource failed"
17     else
18         hs.status = "Progressing"
19         hs.message = "Waiting for resource to be ready"
20     end
21 end
22
23 return hs
```

Custom Lua health check for CRDs that ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) does not know how to assess.

Diff/Drift Detection Issues

Sometimes ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) reports a resource as `OutOfSync` even though you haven't changed anything. You click "Sync," it completes successfully, and then immediately shows `OutOfSync` again. This is the dreaded sync loop, and it usually means something in the cluster is modifying the resource after ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) applies it.

The most common culprit is the Horizontal Pod Autoscaler. Your manifest says `replicas: 3`, but HPA (Horizontal Pod Autoscaler) scales it to 5. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) sees the difference and wants to "fix" it back to 3. HPA (Horizontal Pod Autoscaler) scales it back to 5. Repeat forever.

The fix is to tell ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) to ignore fields that are managed by other controllers:

argocd-apps/myapp.yaml

```
1 # ArgoCD Application manifest with ignoreDifferences.
2 # This goes in your ArgoCD application definition, not
```




```
3  # in the application manifests themselves. Use it to
4  # tell ArgoCD which fields are managed by other controllers.
5
6  apiVersion: argoproj.io/v1alpha1
7  kind: Application
8  metadata:
9    name: myapp
10 spec:
11   ignoreDifferences:
12     # Ignore fields managed by controllers
13     - group: apps
14       kind: Deployment
15       jsonPointers:
16         - /spec/replicas # Managed by HPA
17     # Ignore cluster-specific annotations
18     - group: ""
19       kind: Service
20       jsonPointers:
21         - /metadata/annotations/cloud.google.com~1neg-status
22     # Ignore all managedFields (Kubernetes internal)
23     - group: "*"
24       kind: "*"
25   managedFieldsManagers:
26     - kube-controller-manager
```

Application-level ignoreDifferences configuration for expected drift.

The `ignoreDifferences` configuration uses JSON pointers to target specific fields, so you need to know exactly which path is causing the drift. Start with `argocd app diff myapp` to see what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) considers out of sync, then add the relevant pointer. Be precise – ignoring too broadly (like an entire `metadata` block) can mask real configuration drift that you actually want ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) to catch.

INFO

If a resource shows `OutOfSync` but the diff looks identical, check for whitespace differences, field ordering, or default values that Kubernetes adds but your manifests do not include.



Other common sources of drift:

- **Mutating webhooks**
That add labels or annotations to resources
- **Kubernetes defaulting**
That adds fields your manifest doesn't specify (like `imagePullPolicy: IfNotPresent`)
- **Controller timestamps**
Like `lastTransitionTime` in status fields
- **Cloud provider annotations**
Added by load balancer controllers or ingress controllers

Use `argocd app diff myapp` to see exactly what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) thinks is different. If the diff shows fields you don't care about, add them to `ignoreDifferences`.

Debugging Workflow

Now that you understand *what* can fail, let's talk about *how* to investigate when it does. When a sync fails, resist the urge to immediately start poking at random things. A systematic approach gets you to the root cause faster. I follow a four-step process: identify where in the sync process the failure occurred, check the relevant logs, inspect the actual Kubernetes state, then reproduce and fix.

Step 1: Identify the Failure Point

Start broad and narrow down. The ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) CLI gives you a quick overview of what's happening:

debug-step-1.sh

```
1  #!/bin/bash
2
3  # Get application status overview
4  argocd app get myapp
5
6  # Get detailed sync status
7  argocd app get myapp --show-operation
8
```



```
9  # List resources with their sync/health status
10 argocd app resources myapp
11
12 # Get specific resource status
13 argocd app resources myapp --kind Deployment --name app
```

Commands for identifying where in the sync process failure occurred.

The `argocd app get` output tells you the sync status (Synced, OutOfSync, Unknown) and health status (Healthy, Progressing, Degraded, Missing). If sync status is OutOfSync but health is Healthy, you have a drift problem. If health is Degraded, something failed to start. If health is Progressing for a long time, something is stuck.

The `--show-operation` flag shows the current or last sync operation, including which phase it's in and any error messages. This is where you'll see "Running PreSync hooks" if a hook is stuck, or "Sync error" with a message if apply failed.

Step 2: Examine ArgoCD Logs

Once you know roughly where the failure is, check the logs for the relevant ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) component. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) has several pods, and each handles different parts of the sync process:

debug-step-2.sh

```
1  #!/bin/bash
2
3  # ArgoCD application controller logs (sync decisions)
4  kubectl logs -n argocd -l app.kubernetes.io/name=argocd-application-controller -f
5
6  # ArgoCD repo server logs (manifest generation)
7  kubectl logs -n argocd -l app.kubernetes.io/name=argocd-repo-server -f
8
9  # ArgoCD server logs (API/UI issues)
10 kubectl logs -n argocd -l app.kubernetes.io/name=argocd-server -f
11
12 # Filter for specific application
13 kubectl logs -n argocd -l app.kubernetes.io/name=argocd-application-controller | grep
    "myapp"
```



ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) component logs for different failure types.

Each ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) component owns a different part of the sync pipeline, so knowing which pod to check narrows your search immediately. The application-controller handles sync orchestration and health evaluation; the repo-server handles manifest generation from Git; and the server pod handles API requests and the UI. Here's the mapping:

Component	Log Location	Failure Types
application-controller	Controller pod	Sync decisions, health checks
repo-server	Repo server pod	Manifest generation, Helm/Kustomize
server	Server pod	API errors, authentication
redis	Redis pod	Caching issues, state corruption

ArgoCD components and their log relevance.

Most sync failures show up in the application-controller logs. If your issue is “manifest generation failed” or you’re seeing Helm/Kustomize errors, check the repo-server logs instead. The server pod logs are rarely useful for sync issues – they’re more for API and authentication problems.

Step 3: Inspect Kubernetes State

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)’s view of the world comes from Kubernetes, so when ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) reports a problem, verify what’s actually happening in the cluster. Sometimes ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)’s status is stale, or the problem is a Kubernetes issue that ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) is just reporting.

```
debug-step-3.sh
```

```
1  #!/bin/bash
2
3  # Check deployment rollout status
```



```

4  kubectl rollout status deployment/app -n myapp
5
6  # Describe deployment for events
7  kubectl describe deployment app -n myapp
8
9  # Check pod status and events
10 kubectl get pods -n myapp
11 kubectl describe pod -n myapp -l app=myapp
12
13 # Check recent events in namespace
14 kubectl get events -n myapp --sort-by='.lastTimestamp' | tail -20
15
16 # Check resource quotas if pods pending
17 kubectl describe resourcequota -n myapp

```

Kubernetes commands for inspecting actual cluster state during sync failures.

The `kubectl describe` output for pods is often the most useful. Look at the Events section at the bottom – it shows why pods failed to schedule, failed to pull images, or failed to start containers. Common culprits: image pull failures (wrong tag, missing pull secret), resource quota exceeded, node selector constraints not satisfied.

Step 4: Reproduce and Fix

Once you understand the failure, you need to fix it and verify the fix works. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) provides several options for testing changes before committing them:

debug-step-4.sh

```

1  #!/bin/bash
2
3  # Dry-run sync to see what would happen
4  argocd app sync myapp --dry-run
5
6  # Sync with increased verbosity
7  argocd app sync myapp --loglevel debug
8
9  # Force sync (skip pruning safety)
10 argocd app sync myapp --force
11
12 # Sync specific resources only

```



```
13  argocd app sync myapp --resource apps:Deployment:app
14
15  # Hard refresh (clear cache, re-fetch from Git)
16  argocd app get myapp --hard-refresh
```

Commands for reproducing and testing sync fixes.

The `--dry-run` flag shows you what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) *would* do without actually doing it. Use this to verify your fix before applying it.

The `--resource` flag lets you sync a single resource, which is useful when you’ve fixed one thing and want to test it without re-syncing everything.

The `--hard-refresh` flag clears ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)’s cache and re-fetches manifests from Git. Use this when ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) seems to be showing stale state, or when you’ve made changes to the repo and ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) hasn’t picked them up yet.

Be careful with `--force`. It tells ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) to delete and recreate resources instead of patching them, which can cause downtime. Only use it when you’re stuck and understand the consequences.

Specific Failure Scenarios

The debugging workflow above applies to any failure, but some scenarios come up often enough that they deserve specific runbooks. Here are the four I encounter most frequently.

Scenario: Sync Stuck on Hook

You triggered a sync, and it’s been sitting at “Running PreSync hooks” for 10 minutes. The ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) UI shows a spinner. Nothing seems to be happening. This is one of the most frustrating failure modes because there’s no error message — just silence.

The root cause is almost always a hook Job that isn’t completing. Either it’s still running (legitimately slow), it’s stuck waiting for something (resource contention, external dependency), or it crashed and Kubernetes is respecting the backoff retry delay.



Here's the runbook I use:

runbooks/hook-stuck.md

```

1  # Runbook: Sync Stuck on Hook
2
3  ### Symptoms
4  - Sync status shows "Running PreSync hooks" or "Running PostSync hooks"
5  - Sync has been in this state for longer than expected
6
7  ### Diagnosis
8
9  1. Identify the stuck hook:
10
11  ```bash
12  kubectl get jobs -n myapp -l argocd.argoproj.io/hook
13  ```
14
15  2. Check hook job status:
16
17  ```bash
18  kubectl describe job <hook-job-name> -n myapp
19  ```
20
21  3. Check hook pod logs:
22
23  ```bash
24  kubectl logs -n myapp -l job-name=<hook-job-name>
25  ```
26
27  ### Common Causes
28
29  - Hook script error: Fix script, redeploy
30  - Missing permissions: Add RBAC for hook ServiceAccount
31  - Image pull failure: Check image name and pull secret
32  - Resource limits: Increase memory/CPU limits
33  - External dependency: Ensure dependency is available
34
35  ### Resolution
36
37  1. If hook can be safely skipped:
38
39  ```bash
40  argocd app sync myapp --prune --force

```



```
41  ```
42
43  2. If hook needs to complete:
44
45  - Fix the underlying issue
46  - Delete the stuck job
47  - Trigger new sync
48
49  3. To prevent recurrence:
50
51  - Add `activeDeadlineSeconds` to job spec
52  - Add proper health checks to hook
53  - Consider using `hook-delete-policy: BeforeHookCreation`
```

Runbook for debugging sync stuck on hook execution.


The key insight is that ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) can't tell you why the hook is stuck – it only knows the Job hasn't reached a terminal state. You have to dig into Kubernetes to find the actual cause.

Scenario: OutOfSync Loop

This one is maddening. You sync the application, it completes successfully, and five seconds later it shows OutOfSync again. You sync again. Same result. The application is perpetually out of sync no matter how many times you sync it.

The cause is always something modifying the resource after ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) applies it. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) applies your manifest, a controller or webhook modifies the resource, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) detects the difference and marks it OutOfSync.





OutOfSync loop caused by external modifications to resources

To diagnose, run `argocd app diff myapp` and look at what's different. The diff will show you exactly which fields are changing:

argocd-apps/myapp-with-ignore.yaml

```

1  # Check what ArgoCD sees as different
2  # Run: argocd app diff myapp
3
4  # Example diff output showing HPA modifying replicas:
5  # --- /apps/v1/Deployment/myapp/app
6  # +++ live
7  # @@ -15,7 +15,7 @@
8  #     spec:
  
```



```
9  # -   replicas: 3
10 # +   replicas: 5 # Modified by HPA
11
12 # Solution: Add to Application spec
13 spec:
14   ignoreDifferences:
15     - group: apps
16       kind: Deployment
17       jsonPointers:
18         - /spec/replicas
```

Diagnosing and fixing OutOfSync loop caused by HPA (Horizontal Pod Autoscaler) replica management.

Once you've identified the drifting field, add it to `ignoreDifferences` and sync again. The loop should stop immediately. If it doesn't, re-run the diff – there may be multiple fields drifting, or a mutating webhook might be adding new fields you haven't accounted for yet.

WARNING

Never ignore differences on fields that represent actual drift you care about. Only ignore fields that are legitimately managed by other controllers (HPA (Horizontal Pod Autoscaler) replicas, cert-manager annotations, etc.).

Scenario: CRD Sync Order Problem

Custom Resource Definitions have a chicken-and-egg problem: you can't create a Custom Resource until its CRD (Custom Resource Definition) exists, but if both are in the same ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) application, they might be applied in the wrong order.

The symptom is a sync failure with an error like `the server doesn't have a resource type "MyResource"`. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) tried to create the Custom Resource before the CRD (Custom Resource Definition) was registered with the API server.



The fix has two parts. First, put the CRD (Custom Resource Definition) in an earlier sync wave so it's applied first. Second, add `SkipDryRunOnMissingResource=true` to the CRD (Custom Resource Definition) because ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)'s dry-run will fail if the CRD (Custom Resource Definition) doesn't exist yet:

manifests/crds/myresource-crd.yaml

```
1  # CRD definition - must be in earlier sync wave.
2  # The SkipDryRunOnMissingResource option prevents dry-run failures
3  # when the CRD doesn't exist yet in the cluster.
4
5  apiVersion: apiextensions.k8s.io/v1
6  kind: CustomResourceDefinition
7  metadata:
8    annotations:
9      argocd.argoproj.io/sync-wave: "-5"
10     argocd.argoproj.io/sync-options: SkipDryRunOnMissingResource=true
11  name: myresources.example.com
12  spec:
13    # ... CRD spec
14    ---
15  # CRD instance - must wait for CRD to exist
16  apiVersion: example.com/v1
17  kind: MyResource
18  metadata:
19    annotations:
20      argocd.argoproj.io/sync-wave: "0"
21  name: my-instance
22  spec:
23    # ... instance spec
```

CRD (Custom Resource Definition) and instance ordering with sync waves to prevent race condition.

I use wave -5 for CRDs because I want plenty of room for other infrastructure resources (Namespaces at -3, Secrets at -1, etc.). The specific number doesn't matter as long as it's lower than the wave of any resources that use the CRD (Custom Resource Definition).



Scenario: Resource Pruning Gone Wrong

Pruning is ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)'s way of deleting resources that exist in the cluster but not in Git. It's powerful and dangerous. If you accidentally remove a manifest from Git, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) will helpfully delete the corresponding resource from your cluster. Sometimes that's a Secret you definitely didn't want deleted.

The most common pruning accidents happen when:

You refactor your manifests and forget to include something.

You rename a resource. ArgoCD sees it as *"delete old, create new"*.

A merge conflict resolution drops a file.

You're testing changes in a branch and accidentally merge to main.

Protect critical resources with the `Prune=false` annotation. This tells ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) to never prune this resource, even if it disappears from Git:

manifests/app/secrets/critical-secret.yaml

```
1  # Protect critical resources from accidental pruning.
2  # The Prune=false annotation ensures this resource survives
3  # even if it's accidentally removed from the Git repository.
4
5  apiVersion: v1
6  kind: Secret
7  metadata:
8    annotations:
9      argocd.argoproj.io/sync-options: Prune=false
10   name: critical-secret
11  ---
12  # Application-level prune settings go in your ArgoCD Application manifest
13  apiVersion: argoproj.io/v1alpha1
14  kind: Application
```



```
15 spec:
16   syncPolicy:
17     automated:
18       prune: true
19       selfHeal: true
20   syncOptions:
21     # Don't prune if sync fails
22     - PrunePropagationPolicy=foreground
23     # Require manual prune for certain resources
24     - PruneLast=true
```

Prune protection for critical resources and safe prune policies.

Beyond per-resource protection, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) offers several application-level sync options that control pruning behavior:

#	Sync Option	Effect
1	Prune=false	Never prune this resource
2	PruneLast=true	Prune after all other resources synced
3	PrunePropagationPolicy=foreground	Wait for dependents before pruning
4	ApplyOutOfSyncOnly=true	Only apply resources that differ

ArgoCD sync options and their effects.

Table: ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) sync options affecting prune behavior.

For applications with automated sync enabled, I recommend **PruneLast=true** at the application level. This ensures pruning only happens after all other resources are healthy, reducing the blast radius if something goes wrong during the sync.



Prevention Strategies

Debugging sync failures is valuable, but preventing them is better. The most effective teams I've worked with treat sync failures as signals to improve their deployment process, not just problems to fix. After enough incidents, patterns emerge – and those patterns point to prevention strategies.

The goal isn't to eliminate all sync failures (that's impossible), but to catch the preventable ones early and make the inevitable ones easier to diagnose.

Manifest Validation

Most sync failures I've seen could have been caught before the manifests ever reached ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes). A typo in a resource name, an invalid API version, a missing required field – these are CI problems, not deployment problems. Shift validation left and you'll dramatically reduce sync failures.

The key is layering validation: syntax checking catches malformed YAML (YAML Ain't Markup Language), schema validation catches invalid Kubernetes resources, and dry-run testing catches issues with your actual cluster state.

github-action-validation.yaml

```
1  name: Validate Manifests
2  on: [pull_request]
3
4  jobs:
5    validate:
6      runs-on: ubuntu-latest
7      steps:
8        - uses: actions/checkout@v4
9
10       - name: Install tools
11         run: |
12           curl -LO
13             https://github.com/instrumenta/kubeval/releases/latest/download/kubeval-linux-
14             amd64.tar.gz
15           tar xf kubeval-linux-amd64.tar.gz
16           sudo mv kubeval /usr/local/bin/
```



```
16     - name: Validate YAML syntax
17       run: |
18         find . -name '*.yaml' -exec yamllint {} \;
19
20     - name: Validate Kubernetes schemas
21       run: |
22         kubeval --strict manifests/**/*.yaml
23
24     - name: Dry-run against cluster
25       run: |
26         argocd app diff myapp --local ./manifests --exit-code
```

CI pipeline for manifest validation before merge.

The `argocd app diff --local` step is particularly valuable — it compares what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) *would* generate against what’s currently deployed. This catches not just invalid manifests but manifests that would cause unexpected changes. I’ve seen this single check prevent dozens of incidents where someone accidentally reverted a hotfix or overwrote a manual configuration.

💡 TIP

Run `kubeval` with `--strict` mode to catch unknown fields. Kubernetes accepts unknown fields silently, which means typos in field names (like `replcia` instead of `replica`) won’t cause errors — they’ll just be ignored. Strict mode catches these.

Sync Windows

Not every hour is a good time to deploy. Sync windows let you control *when* ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) can apply changes, which is critical for production environments where you need predictability.

```
argocd-projects/production-sync-windows.yaml
```

```
1  # AppProject sync window configuration.
2  # Applies to all applications within this project.
```



```

3  apiVersion: argoproj.io/v1alpha1
4  kind: AppProject
5  metadata:
6    name: production
7  spec:
8    syncWindows:
9      # Allow syncs only during business hours
10     - kind: allow
11       schedule: "0 9 * * 1-5" # 9 AM weekdays
12       duration: 8h
13       applications:
14         - "*"
15     # Block syncs during peak traffic
16     - kind: deny
17       schedule: "0 12 * * *" # Noon daily
18       duration: 2h
19       applications:
20         - "critical-*"
21     # Allow manual syncs anytime for emergencies
22     - kind: allow
23       schedule: "* * * * *"
24       duration: 24h
25       manualSync: true

```

Sync window configuration restricting automated deployments to safe periods.

The `manualSync: true` window is important – it lets you override the restrictions for emergency deployments. Without this escape hatch, sync windows can become a liability during incidents when you *need* to deploy a fix immediately.

I typically set up three types of windows: an allow window during business hours when the team is available to respond to issues, a deny window during peak traffic when failed deployments have maximum impact, and an always-on manual window for emergencies. The specific times depend on your traffic patterns and team distribution.

Monitoring and Alerting

Sync failures are inevitable. What matters is how quickly you detect and respond to them. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) exposes Prometheus metrics that make it straightforward to alert on sync problems before they become user-facing incidents.



argocd-alerts.yaml

```
1  groups:
2    - name: argocd
3      rules:
4        - alert: ArgoCDSyncFailed
5          expr: |
6            argocd_app_info{sync_status="OutOfSync"} == 1
7            and
8            argocd_app_info{health_status!="Healthy"} == 1
9          for: 10m
10         labels:
11           severity: warning
12         annotations:
13           summary: "ArgoCD app {{ $labels.name }} sync failed"
14           runbook: "https://wiki/runbooks/argocd-sync-failed"
15
16        - alert: ArgoCDSyncStuck
17          expr: |
18            time() - argocd_app_info{sync_status="Syncing"} > 1800
19          labels:
20            severity: critical
21          annotations:
22            summary: "ArgoCD app {{ $labels.name }} sync stuck for 30+ minutes"
23
24        - alert: ArgoCDAppDegraded
25          expr: |
26            argocd_app_info{health_status="Degraded"} == 1
27          for: 5m
28          labels:
29            severity: critical
30          annotations:
31            summary: "ArgoCD app {{ $labels.name }} is degraded"
```

Prometheus alerting rules for ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) sync failures.

The three alerts above cover the most common failure modes: sync failures (OutOfSync + unhealthy), stuck syncs (syncing for too long), and degraded applications (healthy but degraded status). The `for` clause prevents alert noise from transient states – a brief OutOfSync during a normal deployment shouldn't page anyone.



✓ SUCCESS

Alert on sync failures early. A 10-minute OutOfSync alert gives you time to investigate before users notice. A 30-minute stuck sync is almost always a real problem requiring intervention.

The runbook link in the annotations is critical. When an engineer gets paged at 3 AM, they shouldn't have to remember the debugging workflow – they should be able to follow a documented procedure. Link each alert to a runbook that walks through the investigation steps.

Advanced Troubleshooting

When the standard debugging workflow doesn't reveal the problem, you need to dig deeper into ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)'s internals. These techniques address edge cases that arise in complex deployments or after upgrades.

Resource Tracking Methods

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) needs to know which resources belong to which application. It tracks this ownership through labels or annotations on the resources themselves. Most of the time, the default label-based tracking works fine – but it can cause problems when those labels interfere with other systems.

The most common issue I've seen is with label selectors. If your application uses label selectors that happen to match ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)'s tracking labels, you can get unexpected behavior. Switching to annotation-based tracking resolves this, since annotations aren't used in selectors.

The table below summarizes when to use each tracking method:

#	Method	How It Works	When to Use
1	label	Adds ArgoCD label to resources	Default, simple apps
2	annotation	Adds ArgoCD annotation	When labels cause issues



#	Method	How It Works	When to Use
3	annotation+label	Both	Migration between methods

ArgoCD resource identification methods.

To change the tracking method or enable related sync options, configure them in your Application spec. The example below shows server-side apply, which is particularly useful when combined with annotation tracking for resources managed by multiple controllers:

argocd-apps/myapp-server-side-apply.yaml

```

1  apiVersion: argoproj.io/v1alpha1
2  kind: Application
3  metadata:
4    name: myapp
5  spec:
6    syncPolicy:
7      syncOptions:
8        # Use annotation tracking if labels interfere with selectors
9        - ServerSideApply=true
10       - RespectIgnoreDifferences=true

```

Application configuration using server-side apply for complex resources.

Server-side apply (`ServerSideApply=true`) is worth calling out here. It delegates conflict resolution to the Kubernetes API server rather than ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes), which handles field ownership more gracefully. If you're seeing persistent drift on resources managed by multiple controllers (like HPA (Horizontal Pod Autoscaler) modifying replica counts), server-side apply often resolves the conflict.

Debugging Manifest Generation

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) doesn't deploy your YAML (YAML Ain't Markup Language) files directly – it renders them through Helm, Kustomize, or plain directory processing first. When sync fails with manifest-related errors, you need to see what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) actually generated, which might differ from what you expect.



The most common surprise is values resolution in Helm. ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) might be using different values files than you're testing locally, or environment-specific overrides might not be applied in the order you expect.

debug-manifest-generation.sh

```

1  #!/bin/bash
2
3  # See what ArgoCD generates from your repo
4  argocd app manifests myapp
5
6  # Compare local vs remote generation
7  argocd app diff myapp --local ./path/to/manifests
8
9  # For Helm apps, check values resolution
10 argocd app get myapp -o yaml | grep -A 50 'helm:'
11
12 # Test Helm template locally
13 helm template myapp ./chart -f values.yaml -f values-prod.yaml
14
15 # Test Kustomize locally
16 kustomize build ./overlays/production
17
18 # Check repo-server can access the repo
19 kubectl exec -n argocd deploy/argocd-repo-server -- \
20   ls /tmp/<repo-hash>/

```

Commands for debugging manifest generation from Helm and Kustomize.

The `argocd app manifests` command is your friend here – it shows you exactly what ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) will apply, after all the templating. Compare this output to what you generate locally, and discrepancies will jump out. The repo-server access check at the end catches a common issue where ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) can't clone or access your repository due to network or authentication problems.

Recovering from Corrupt State

Sometimes ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)'s internal state becomes inconsistent with reality. This can happen after failed upgrades, network partitions, or manual resource modifications that confuse the controller. The symptoms vary – perpetual OutOfSync, missing resources in the UI, or sync operations that silently do nothing.



recovery-commands.sh

```

1  #!/bin/bash
2
3  # Step 1: Try this first - force refresh from Git
4  argocd app get myapp --hard-refresh
5
6  # Step 2: If refresh doesn't help - delete and recreate app (preserves resources)
7  argocd app delete myapp --cascade=false
8  argocd app create myapp --repo ... --path ... --dest-server ...
9
10 # Step 3: If state is still corrupt - clear Redis cache
11 kubectl delete pod -n argocd -l app.kubernetes.io/name=argocd-redis
12
13 # Step 4: Nuclear option - restart application controller
14 kubectl rollout restart deployment -n argocd argocd-application-controller

```

Recovery commands for ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) state corruption scenarios.

Start with `--hard-refresh`, which forces ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) to re-read the Git repository and recalculate the desired state. This resolves most caching issues. If that doesn't work, the nuclear option is deleting and recreating the Application resource – but *only* with `--cascade=false`, which preserves the actual Kubernetes resources while resetting ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes)'s tracking state.

⚠ DANGER

The `--cascade=false` flag is critical when deleting an application for recovery. Without it, ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) will delete all the Kubernetes resources the application manages.

Conclusion

ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) sync failures are frustrating because they break the promise of GitOps (Git as single source of truth for declarative infrastructure): you pushed to Git, so it should just work. But that frustration fades once you understand what's actually happening beneath the



abstraction.

The debugging workflow follows a clear pattern. Start with the ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) UI or CLI to understand *what* failed. Check sync waves and hooks to understand *when* it failed. Examine Kubernetes state to understand *why* it failed. And trace the resource dependency graph to understand whether the failure is isolated or will cascade.

Most sync failures fall into a few categories: resource dependency issues (usually ordering problems), hook failures (scripts that time out or crash), health check failures (resources that don't become ready), and drift detection issues (differences between desired and live state that shouldn't exist). Once you recognize the category, the fix is usually straightforward.

✓ SUCCESS

GitOps (Git as single source of truth for declarative infrastructure) is not magic — it is automation. When the automation fails, you need to understand what it was trying to do. Master the sync process, and debugging becomes systematic rather than frustrating.

The prevention strategies matter as much as the debugging skills. CI validation catches manifest errors before they reach ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes). Sync windows prevent deployments during high-risk periods. Alerting ensures you know about failures before users do. And runbooks mean that when things fail at 3 AM, you have a documented path to resolution.

-
- 1 Lua is the only scripting language supported by ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) for custom health assessments and resource actions. and tell ArgoCD (Declarative GitOps continuous delivery tool for Kubernetes) how to interpret your CRD (Custom Resource Definition)'s status field: ↩



Copyright © 2022 Webstack Builders, Inc.

The text, diagrams, and images in this work are licensed under CC BY-NC 4.0

All code samples in this article are licensed under the MIT License. Feel free to use, modify, and distribute them in any project.

<https://www.webstackbuilders.com/articles/argocd-sync-failures-gitops-debugging-troubleshooting>

