

Ingress vs Gateway API: A Practical Comparison



Published on January 7, 2024



Webstack
Builders

Table of Contents

Understanding the Models	3
Ingress Model	3
Gateway API Model	5
Feature Comparison	8
Routing Capabilities	8
Traffic Management	10
Protocol Support	12
Role-Based Ownership	13
Multi-Team Architecture	14
Migration Strategy	16
Migration Readiness Checklist	17
Incremental Migration Path	18
Ingress to HTTPRoute Conversion	19
Ecosystem Maturity	21
Controller Support	21
Tooling and Documentation	22
Decision Framework	23
When to Choose Ingress	23
When to Choose Gateway API	24
The Hybrid Approach	25
Conclusion	25



Kubernetes Ingress has been the standard for HTTP routing since 2015, but its limitations led to every ingress controller implementing different annotations and behaviors. If you've ever tried to port an Ingress configuration from NGINX to Traefik, you've felt this pain – annotations that worked perfectly on one controller are silently ignored by another. Gateway API, which launched in 2020 and reached GA (General Availability) in 2023, addresses these limitations with a more expressive, portable, and role-oriented design.

But more powerful doesn't always mean better for your use case.

I've seen this play out both ways. A team running NGINX Ingress Controller decided to migrate to Gateway API because "it's the future." They spent two weeks converting resources, only to discover that their simple host-based routing with TLS (Transport Layer Security) termination worked perfectly with Ingress – Gateway API added complexity without new capabilities. Meanwhile, another team I worked with was managing a multi-tenant platform with complex traffic splitting and cross-namespace routing. They'd accumulated a mess of hacky Ingress annotations to work around limitations. For them, Gateway API solved real problems and simplified their configuration.

INFO

There's no deprecation timeline for Ingress. The Kubernetes networking SIG has stated both APIs will be supported indefinitely. Plan your technology choices assuming Ingress remains viable for at least the next 5-10 years.

The lesson: evaluate based on your actual requirements, not industry hype. The choice between Ingress and Gateway API depends on your current tooling, team structure, and routing complexity – not on which technology is newer.

Understanding the Models

Before diving into features and migration, it helps to understand how these two APIs think about routing. Ingress is flat – one resource type handles everything. Gateway API is layered – different resource types for different concerns, managed by different teams.



Ingress Model

Ingress uses a single resource type that combines routing rules, TLS (Transport Layer Security) configuration, and backend references. Everything lives in one YAML file, which makes simple cases easy but complex cases awkward. The catch is that Ingress itself only defines the *interface* – the actual behavior depends entirely on which controller you're running. NGINX Ingress Controller interprets annotations one way; Traefik interprets them differently; some controllers ignore certain annotations entirely.

ingress-example.yaml

```
1  # Single Ingress resource handles routing, TLS, and backend config
2  apiVersion: networking.k8s.io/v1
3  kind: Ingress
4  metadata:
5    name: web-ingress
6    namespace: production
7    annotations:
8      # Controller-specific configuration (not portable)
9      nginx.ingress.kubernetes.io/rewrite-target: /
10     nginx.ingress.kubernetes.io/ssl-redirect: "true"
11     nginx.ingress.kubernetes.io/proxy-body-size: "50m"
12  spec:
13    ingressClassName: nginx
14    tls:
15      - hosts:
16          - api.example.com
17        secretName: api-tls-secret
18    rules:
19      - host: api.example.com
20        http:
21          paths:
22            - path: /v1
23              pathType: Prefix
24              backend:
25                service:
26                  name: api-v1
27                  port:
28                    number: 80
29            - path: /v2
30              pathType: Prefix
```



```
31         backend:
32             service:
33                 name: api-v2
34                 port:
35                     number: 80
```

Ingress resource combining routing, TLS (Transport Layer Security), and backends in one resource.

This simplicity has real limitations. Ingress only supports HTTP and HTTPS – no TCP (Transmission Control Protocol), UDP (User Datagram Protocol), or native gRPC (Google Remote Procedure Call) routing. Configuration portability is poor because any non-trivial behavior requires controller-specific annotations. All resources must live in the same namespace as the services they route to. And there's no separation between infrastructure concerns (which controller, what TLS (Transport Layer Security) policy) and application concerns (which paths route where).

Gateway API Model

Gateway API takes a different approach: separate resource types for separate concerns, with explicit ownership boundaries between infrastructure teams, platform teams, and application teams.

The architecture has three layers. **GatewayClass** is cluster-scoped and defines which controller implementation handles Gateways – typically managed by the infrastructure team. **Gateway** is namespace-scoped and defines listeners (ports, protocols, TLS (Transport Layer Security) termination) – typically managed by the platform team. **HTTPRoute** (and other Route types) defines the actual routing rules – typically managed by application teams who only need to reference the Gateway they want to attach to.

gateway-api-example.yaml

```
1  # Layer 1: GatewayClass (cluster-scoped, managed by infra team)
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: GatewayClass
4  metadata:
5      name: production-gateway-class
6  spec:
7      controllerName: gateway.nginx.org/nginx-gateway-controller
8      description: "Production NGINX Gateway"
9
```



```

10  ---
11  # Layer 2: Gateway (namespace-scoped, managed by platform team)
12  apiVersion: gateway.networking.k8s.io/v1
13  kind: Gateway
14  metadata:
15    name: production-gateway
16    namespace: gateway-infra
17  spec:
18    gatewayClassName: production-gateway-class
19    listeners:
20      - name: https
21        port: 443
22        protocol: HTTPS
23        hostname: "*.example.com"
24        tls:
25          mode: Terminate
26          certificateRefs:
27            - name: wildcard-tls
28              namespace: gateway-infra
29        allowedRoutes:
30          namespaces:
31            from: Selector
32            selector:
33              matchLabels:
34                gateway-access: "true"
35  ---
36  # Layer 3: HTTPRoute (namespace-scoped, managed by app teams)
37  apiVersion: gateway.networking.k8s.io/v1
38  kind: HTTPRoute
39  metadata:
40    name: api-routes
41    namespace: production
42  spec:
43    parentRefs:
44      - name: production-gateway
45        namespace: gateway-infra # Cross-namespace reference
46    hostnames:
47      - "api.example.com"
48    rules:
49      - matches:
50        - path:

```




```

52         type: PathPrefix
53         value: /v1
54     backendRefs:
55     - name: api-v1
56       port: 80
57   - matches:
58     - path:
59         type: PathPrefix
60         value: /v2
61     backendRefs:
62     - name: api-v2
63       port: 80

```

Gateway API's layered model separates infrastructure, platform, and application concerns.

Notice how the HTTPRoute in the `production` namespace can reference a Gateway in the `gateway-infra` namespace. This cross-namespace capability is fundamental to Gateway API's multi-team model – application teams don't need access to the Gateway namespace to define their routes.



Ingress uses flat resources - Gateway API uses a shared Gateway.



The table below summarizes the key differences. If you're scanning for a quick answer: Ingress is simpler when you don't need the advanced features, Gateway API is more capable when you do.

#	Aspect	Ingress	Gateway API
1	Resource types	1 (Ingress)	3+ (GatewayClass, Gateway, *Route)
2	Configuration	Annotations (not portable)	Native fields (portable)
3	Namespace scope	Same namespace only	Cross-namespace supported
4	Protocols	HTTP/HTTPS	HTTP, HTTPS, TCP, UDP, gRPC, TLS
5	Traffic splitting	Annotation-dependent	Native support
6	Header manipulation	Annotation-dependent	Native filters
7	Role separation	No	Yes (infra/platform/app)

Ingress vs Gateway API comparison.

⚠ WARNING

Gateway API's layered model adds complexity. If you don't need cross-namespace routing, traffic splitting, or role separation, the extra layers are overhead without benefit. Ingress's simplicity is a feature, not a limitation.

Feature Comparison

Routing Capabilities

Beyond basic host and path matching, modern routing often requires matching on headers, query parameters, or HTTP methods. Ingress provides limited support for these patterns – you're dependent on controller-specific annotations that may or may not exist. Gateway API makes advanced matching a first-class feature.



#	Capability	Ingress	Gateway API
1	Path matching	Exact, Prefix, ImplementationSpecific	Exact, PathPrefix, RegularExpression
2	Header matching	Annotation-dependent	Native (`matches.headers`)
3	Query parameter matching	Not supported	Native (`matches.queryParams`)
4	HTTP method matching	Not supported	Native (`matches.method`)

Routing capability comparison.

The difference becomes concrete when you need to route requests based on something other than URL path. With Gateway API, routing beta users to a different backend based on a header is straightforward:

header-based-routing.yaml

```

1  # Gateway API: Route beta users to canary backend
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: HTTPRoute
4  metadata:
5    name: api-beta-route
6  spec:
7    parentRefs:
8      - name: production-gateway
9    hostnames:
10     - api.example.com
11    rules:
12     - matches:
13       - headers:
14         - name: X-Beta-User
15           value: "true"
16       path:
17         type: PathPrefix
18         value: /api
19     backendRefs:

```



```
20         - name: api-canary
21           port: 80
```

Header-based routing in Gateway API.

With Ingress, you'd need to check whether your controller supports a custom annotation for header routing – and if you ever switch controllers, that annotation likely won't transfer.

Traffic Management

Traffic splitting for canary deployments is where the annotation problem really shows. With NGINX Ingress Controller, you create a separate “canary” Ingress resource with specific annotations. With Traefik, you use different annotations. With HAProxy, different still. None of these approaches are portable.

Gateway API handles traffic splitting natively through the `weight` field on backend references:

traffic-splitting.yaml

```
1  # Gateway API: Native weighted traffic splitting
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: HTTPRoute
4  metadata:
5    name: api-canary-route
6  spec:
7    parentRefs:
8      - name: production-gateway
9    hostnames:
10     - api.example.com
11    rules:
12     - backendRefs:
13       - name: api-stable
14         port: 80
15         weight: 80
16       - name: api-canary
17         port: 80
18         weight: 20
```

Weighted traffic splitting in Gateway API – no annotations required.



Header manipulation follows the same pattern. Gateway API provides native filters for adding, modifying, or removing request and response headers. URL rewriting is similarly built-in. These capabilities that required controller-specific configuration in Ingress are portable across any Gateway API implementation.

header-manipulation.yaml

```
1  # Gateway API: Request and response header modification
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: HTTPRoute
4  metadata:
5    name: api-with-headers
6  spec:
7    parentRefs:
8      - name: production-gateway
9    rules:
10     - matches:
11       - path:
12         type: PathPrefix
13         value: /api
14     filters:
15       - type: RequestHeaderModifier
16         requestHeaderModifier:
17           add:
18             - name: X-Request-Source
19               value: "gateway"
20           remove:
21             - X-Internal-Only
22       - type: ResponseHeaderModifier
23         responseHeaderModifier:
24           add:
25             - name: X-Served-By
26               value: "api-cluster"
27     backendRefs:
28       - name: api-service
29       port: 80
```

Header manipulation using Gateway API filters.



✓ SUCCESS

Gateway API's native traffic splitting and header manipulation are its strongest advantages. If you're currently using complex Ingress annotations for canary deployments or A/B testing, Gateway API makes these patterns first-class citizens.

Protocol Support

Ingress was designed for HTTP and HTTPS. Period. If you need to route TCP (Transmission Control Protocol) traffic to a database, UDP (User Datagram Protocol) traffic to a DNS server, or gRPC (Google Remote Procedure Call) traffic with method-level routing, Ingress doesn't help. Some controllers added annotation-based workarounds, but these are non-standard and limited.

Gateway API includes dedicated route types for different protocols. TCPRoute and UDPRoute (currently in beta and alpha respectively) enable layer-4 routing. GRPCRoute provides native gRPC (Google Remote Procedure Call) support with service and method matching. TLSRoute handles TLS (Transport Layer Security) passthrough without termination.

grpc-routing.yaml

```
1  # Gateway API: Native gRPC routing with method matching
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: GRPCRoute
4  metadata:
5    name: user-service-route
6  spec:
7    parentRefs:
8      - name: production-gateway
9    hostnames:
10     - grpc.example.com
11    rules:
12     - matches:
13       - method:
14           service: myapp.UserService
15           method: GetUser
16     backendRefs:
```



```
17     - name: user-service
18       port: 9090
```

gRPC (Google Remote Procedure Call) routing with method-level matching in Gateway API.

The table below summarizes protocol support across both APIs. For HTTP workloads, either option works. For anything else, Gateway API is your only portable choice.

Protocol	Ingress	Gateway API	Status
HTTP	✓ Native	✓ HTTPRoute	GA
HTTPS (terminate)	✓ Native	✓ HTTPRoute + TLS	GA
HTTPS (passthrough)	Annotation-dependent	✓ TLSRoute	GA
TCP	Not supported	✓ TCPRoute	Beta
UDP	Not supported	✓ UDPRoute	Alpha
gRPC	Annotation-dependent	✓ GRPCRoute	GA

Protocol support comparison. Status column reflects Gateway API specification maturity (GA/Beta/Alpha), not individual controller implementation status.

Role-Based Ownership

In organizations with multiple teams, the question of *who manages what* becomes critical. With Ingress, there's no natural separation – anyone with permission to create Ingress resources in a namespace can affect cluster-wide routing behavior. This leads to either overly permissive access (risky) or centralized bottlenecks where a platform team manages all routing (slow).

Gateway API's layered architecture maps naturally to organizational boundaries. Each layer corresponds to a different team with different responsibilities and different access levels.



Multi-Team Architecture

The three-layer model creates clear ownership:

➤ **Infrastructure team**

Manages GatewayClass resources. They select and configure the controller, set cluster-wide policies, and typically have cluster-admin access. This layer changes rarely – maybe during major upgrades or when adding new controller capabilities.

➤ **Platform team**

Manages Gateway resources. They configure listeners, TLS certificates, and crucially, which namespaces can attach routes. They control the blast radius of application team changes through the `allowedRoutes` configuration.

➤ **Application teams**

Manage HTTPRoute resources in their own namespaces. They define their routing rules, but can only attach to Gateways that explicitly allow their namespace. They can't affect other teams' traffic.

multi-team-gateway.yaml

```
1  # Platform team's Gateway with namespace restrictions
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: Gateway
4  metadata:
5    name: shared-gateway
6    namespace: gateway-system
7  spec:
8    gatewayClassName: production
9    listeners:
10     - name: https
11       port: 443
12       protocol: HTTPS
13       hostname: "*.example.com"
14       tls:
15         mode: Terminate
16         certificateRefs:
17           - name: wildcard-cert
18       allowedRoutes:
19         namespaces:
20           from: Selector
21           selector:
```




```
22         matchLabels:
23             gateway-access: "true"
24         kinds:
25             - kind: HTTPRoute
26
27     ---
28     # Application team's HTTPRoute (in their namespace)
29     apiVersion: gateway.networking.k8s.io/v1
30     kind: HTTPRoute
31     metadata:
32         name: team-a-routes
33         namespace: team-a # Must have label: gateway-access: "true"
34     spec:
35         parentRefs:
36             - name: shared-gateway
37               namespace: gateway-system
38         hostnames:
39             - "team-a.example.com"
40         rules:
41             - matches:
42                 - path:
43                     type: PathPrefix
44                     value: /
45                 backendRefs:
46                     - name: team-a-service
47                       port: 80
```

Gateway with namespace selector controlling which teams can attach routes.

The `allowedRoutes.namespaces.selector` is the key mechanism. Only namespaces with the label `gateway-access: "true"` can create HTTPRoutes that attach to this Gateway. The platform team controls which namespaces get that label, effectively controlling which teams can self-service their routing.

This separation is enforced at the Kubernetes API level, not just by convention. If an application team tries to create an HTTPRoute referencing a Gateway that doesn't allow their namespace, the route won't attach – it will show an error in its status conditions.





Role-based ownership with Gateway API.

You'll want to back this up with Kubernetes RBAC (Role-Based Access Control). Give the infrastructure team permissions on `GatewayClass` resources, the platform team permissions on `Gateway` resources in the `gateway-system` namespace, and application teams permissions only on `HTTPRoute` resources in their own namespaces. The Gateway API resources and RBAC (Role-Based Access Control) work together to create a self-service model with appropriate guardrails.

INFO

Role separation is Gateway API's killer feature for large organizations. Platform teams can manage shared infrastructure while application teams self-service their routing – with guardrails enforced by the Gateway's `allowedRoutes` configuration.



Migration Strategy

If you've decided Gateway API is right for your use case, you don't need to migrate everything at once. Ingress and Gateway API can coexist on the same cluster, even with the same controller. This enables an incremental approach that minimizes risk.

Migration Readiness Checklist

Assuming you've already decided Gateway API is the right fit (see the Decision Framework section), the next question is whether now is the right time to execute the migration.

Timing signals that suggest waiting:

- ☐ You're in the middle of a major release or incident stabilization period
- ☐ The team responsible for routing config doesn't have capacity for learning and testing
- ☐ Your controller's Gateway API support is still in beta or missing route types you need
- ☐ You have a complex configuration-snippet annotation that doesn't have a clear Gateway API equivalent
- ☐ Your CI / CD pipeline doesn't have good staging coverage for routing changes

Prerequisite checklist before starting:

- ☐ Controller supports Gateway API at GA level for the route types you need — most major controllers (NGINX Gateway Fabric, Envoy Gateway, Istio, Contour, Kong, Traefik) have GA support
- ☐ Team has read the official Gateway API getting started guide
- ☐ Staging environment can run parallel Ingress and Gateway API resources
- ☐ Rollback plan documented (typically: delete HTTPRoute, traffic returns to Ingress)
- ☐ Monitoring covers both Ingress and Gateway API traffic paths
- ☐ At least one low-risk service identified for initial migration



Effort estimation:

Migration timelines depend heavily on annotation complexity. If your Ingress resources use only standard fields (hosts, paths, TLS (Transport Layer Security)), conversion is mechanical and you can migrate a route in under an hour. If you're relying on custom annotations, rate limiting policies, or `configuration-snippet` blocks, plan for research time to find Gateway API equivalents. A typical 20-service cluster with moderate annotation usage takes 4-8 weeks for a careful migration with parallel operation.

Incremental Migration Path

The safest migration runs both APIs in parallel, migrating routes one at a time with the ability to roll back each change independently.

1

Phase 1: Parallel operation (2-4 weeks)

Deploy the GatewayClass and Gateway resources, but only create HTTPRoutes for new services. Your existing Ingress resources continue to handle existing traffic. This phase validates that your Gateway setup works without risking existing routes.

parallel-operation.yaml

```
1  # New API version uses Gateway API
2  apiVersion: gateway.networking.k8s.io/v1
3  kind: HTTPRoute
4  metadata:
5    name: api-v2-route
6    namespace: production
7  spec:
8    parentRefs:
9      - name: production-gateway
10        namespace: gateway-system
11    hostnames:
12      - "api.example.com"
13    rules:
14      - matches:
15          - path:
16              type: PathPrefix
17              value: /v2
```



```
18     backendRefs:
19       - name: api-v2
20         port: 80
21   # Existing /v1 routes still handled by Ingress
```

New routes use Gateway API while existing Ingress handles legacy paths.

This coexistence works because Ingress and Gateway API are separate resource types that controllers watch independently. When a request arrives, the controller checks both APIs and routes based on the matching rules. Requests to `/v2` match the HTTPRoute; requests to `/v1` match the Ingress. There's no conflict because they handle different paths. If you need to migrate an existing path, create the HTTPRoute first with a different hostname for testing, then update DNS or the hostname field to cut over.

2**Phase 2: Migrate routes (4-8 weeks)**

Convert Ingress rules to HTTPRoutes one at a time. Test each migration in staging first, then deploy to production. Keep the original Ingress rule as a fallback — you can revert instantly if something breaks.

3**Phase 3: Cutover (1-2 weeks)**

Once all routes are migrated, route all traffic through Gateway. Keep Ingress resources in read-only mode (no new rules) while you monitor for issues. This is your last chance to catch problems before cleanup.

4**Phase 4: Cleanup (1 week)**

Delete unused Ingress resources, document the new patterns, and update team training. This is the point of no return — make sure you're confident before proceeding.

The key insight is that rollback is easy in phases 1-3. If an HTTPRoute causes problems, delete it and the traffic falls back to Ingress. Only in phase 4, when you delete Ingress resources, do you commit to Gateway API.

Ingress to HTTPRoute Conversion

Converting an Ingress resource to an HTTPRoute is mostly mechanical, but there are gotchas. The path matching types have different names: Ingress `Prefix` becomes HTTPRoute `PathPrefix`. Ingress `ImplementationSpecific` has no direct equivalent — you'll need to understand what your controller actually did with it and pick the appropriate HTTPRoute match type.



Ingress pathType	HTTPRoute path.type	Notes
Exact	Exact	Direct mapping
Prefix	PathPrefix	Direct mapping
ImplementationSpecific	Varies	Check controller docs

Path type mapping between Ingress and HTTPRoute.

The harder part is annotations. Each annotation that your Ingress uses for non-standard behavior needs to be converted to Gateway API's native features or policy attachments. Common conversions:

Rewrite annotations: | URLRewrite filter

Header add/remove annotations: | RequestHeaderModifier / ResponseHeaderModifier filters

Rate limiting annotations: | Policy attachment (controller-specific)

Authentication annotations: | Policy attachment or external auth service

The `configuration-snippet` annotation is the hardest to migrate. It allows arbitrary NGINX config, which means you need to understand exactly what it's doing and find the Gateway API equivalent. Sometimes there isn't one – you may need a controller-specific extension or policy.

⚠ WARNING

Don't migrate during high-traffic periods or with tight deadlines. The parallel operation phase should run long enough to catch edge cases. Rushed migrations create incidents – plan for a multi-week timeline.



Ecosystem Maturity

Gateway API is newer, which means less community content and less battle-tested operational knowledge. That gap is closing, but it's worth understanding where you might hit friction.

Controller Support

All major ingress controllers now support Gateway API at some level. The implementations vary in completeness:

Controller	Gateway API Status	Route Types	Notes
Envoy Gateway	GA	HTTP, gRPC, TCP, UDP, TLS	Purpose-built for Gateway API
Istio	GA	HTTP, gRPC, TCP, TLS	Preferred over Istio Gateway
Contour	GA	HTTP, gRPC, TLS	Mature implementation
Kong	GA	HTTP, gRPC, TCP, UDP	Full plugin support via policy
Traefik	GA	HTTP, gRPC, TCP, TLS	Middleware via policy
NGINX Gateway Fabric	GA	HTTP, gRPC	Successor to nginx-ingress
HAProxy	Beta	HTTP, TCP	Newer, less mature

Controller Gateway API support.

If you're currently using NGINX Ingress Controller, note that Gateway API support comes through a different project – NGINX Gateway Fabric – not the original nginx-ingress. They can coexist, but they're separate deployments.



Tooling and Documentation

Ingress has an 8-year head start. You'll find more blog posts, more examples, and more answers to edge-case questions. Gateway API's official documentation is comprehensive, but community content is still catching up.

Debugging tools are similar in capability but different in approach. Gateway API resources have richer status conditions that tell you exactly why a route isn't working. Run `kubectl describe httproute <name>` and look at the `conditions` section:

Bash

```
1  Status:
2    Parents:
3    Conditions:
4      Type:    Accepted
5      Status:  False
6      Reason:  NotAllowedByListeners
7      Message: Route hostname "api.example.com" does not match listener
```

HTTPRoute status conditions showing why a route failed to attach.

This structured feedback beats digging through controller logs – you can see immediately that the hostname doesn't match the Gateway's listener configuration. Ingress debugging typically means reading controller logs to understand why traffic isn't routing as expected.

For GitOps workflows (ArgoCD, Flux), both work natively. The main gotcha with Gateway API is that you need the CRDs installed before ArgoCD tries to sync resources that use them – plan your sync waves accordingly.

#	Aspect	Ingress	Gateway API
1	Years in production	8+ years	3+ years
2	Stack Overflow questions	10,000+	500+
3	Battle-tested patterns	Extensive	Growing
4	Breaking changes risk	Low (stable API)	Low (v1 is GA)



#	Aspect	Ingress	Gateway API
5	Controller coverage	Universal	~80% of major controllers

Ecosystem maturity comparison.

 INFO

Gateway API reached GA (General Availability) (v1.0) in late 2023. The core APIs (GatewayClass, Gateway, HTTPRoute, GRPCRoute) are stable. Experimental APIs (TCPRoute, UDPRoute) may change. For production HTTP routing, Gateway API is now mature enough for most use cases.

Decision Framework

After all the feature comparisons and migration details, the decision comes down to a few key questions.

When to Choose Ingress

Stick with Ingress if:

Your routing needs are simple: host-based routing, path prefixes, TLS termination	
Everything lives in one namespace (or a small number of namespaces managed by one team)	
Your current setup works and you don't have concrete problems to solve	
Your team doesn't have capacity to learn a new API right now	



Your controller doesn't support Gateway API yet



Example scenarios: a single application with a few endpoints, internal tools with basic routing, straightforward microservices where one team owns everything.

When to Choose Gateway API

Choose Gateway API if:

You need cross-namespace routing (app teams routing to a shared Gateway)	
You need native traffic splitting for canary or blue-green deployments	
You need TCP, UDP, or gRPC routing	
You're building a multi-team platform with self-service routing	
You're planning a controller migration anyway	
It's a greenfield project and you want to start with the more portable option	

Example scenarios: multi-tenant platforms, deployment pipelines with canary stages, mixed-protocol workloads, service mesh integration.

The Hybrid Approach

You don't have to choose one or the other. Many organizations run both:



Ingress for legacy applications and simple internal services where migration effort isn't justified	
Gateway API for new platform services and anything requiring advanced features	

Consider a concrete example: your cluster runs an internal admin dashboard that hasn't changed in two years (Ingress works fine), a legacy API that's in maintenance mode (no reason to migrate), and a new public API with canary deployment requirements (Gateway API). Running both APIs lets you use Gateway API where it adds value without touching stable workloads that don't need its features.

This works because both APIs can coexist on the same cluster, even with the same controller. The main cost is cognitive load – your team needs to know both patterns. Document which services use which API and establish a convention for new services.

Conclusion

The choice between Ingress and Gateway API isn't about which technology is “better” – it's about which fits your current needs and team capabilities.

Ingress is simpler, more mature, and has years of battle-tested operational patterns. If your routing needs are straightforward – host-based routing, path prefixes, TLS (Transport Layer Security) termination in a single namespace – Ingress does the job with less complexity.

Gateway API is more powerful and more portable. Its layered architecture enables multi-team self-service with guardrails. Native support for traffic splitting, header manipulation, and non-HTTP protocols eliminates the annotation sprawl that makes Ingress configurations fragile. If you need these capabilities, Gateway API delivers them in a standardized way that works across controllers.



✓ SUCCESS

The best choice is the one that solves your actual problems with acceptable complexity. If Ingress works for you, keep using it. If you need Gateway API's capabilities, it's mature enough for production. Make the decision based on requirements, not hype.

Both APIs will coexist for years – the Kubernetes networking SIG has confirmed there's no deprecation timeline for Ingress. Migrating purely because Gateway API is “the future” is a poor use of engineering time. Migrate when you have concrete requirements that Gateway API solves better: cross-namespace routing, native traffic splitting, protocol support, or role separation. For greenfield projects where you don't have strong opinions, Gateway API is a reasonable default – it's mature enough for production, and starting there avoids a future migration.

Copyright © 2024 Webstack Builders, Inc.

The text, diagrams, and images in this work are licensed under CC BY-NC 4.0

All code samples in this article are licensed under the MIT License. Feel free to use, modify, and distribute them in any project.

<https://www.webstackbuilders.com/articles/kubernetes-ingress-gateway-api-comparison-migration>

