


OpenAPI in Practice: Docs, Clients, and Validation



Published on December 3, 2023



Webstack
Builders

Table of Contents

Introduction	3
OpenAPI Fundamentals	3
The Building Blocks	3
Schema Design That Doesn't Fight You	5
Composition Patterns	7
Spec-First vs Code-First	9
The Case for Spec-First	10
The Case for Code-First	11
Keeping Spec and Code in Sync	12
Documentation Generation	13
Choosing a Documentation Tool	13
Making Documentation Worth Reading	14
SDK Generation	16
Picking a Generator	16
Generating Your First SDK	18
Making Generated Code Usable	19
Request Validation	20
Adding Validation Middleware	20
Handling Validation Errors	22
What Validation Catches (and Doesn't)	23
CI/CD Integration	24
The Core Pipeline	24
Linting with Spectral	26
Detecting Breaking Changes	27
Contract Testing	27
Conclusion	28



Introduction

OpenAPI promises a single specification file that generates accurate documentation, type-safe client SDKs, and request validation middleware. Write once, generate everywhere. The reality is more nuanced – specs drift from implementations, generated code can be awkward, validation catches schema errors but misses business logic. The difference between OpenAPI delivering value and becoming another artifact to maintain comes down to one question: is the spec the source of truth, or is it describing code that already exists?

I've seen teams maintain REST APIs with hand-written documentation in Confluence, manually coded client SDKs in three languages, and validation logic duplicated across services. Documentation is perpetually out of date. SDK (Software Development Kit) bugs emerge when the API changes. Validation differs between the gateway and the backend. When they adopt spec-first OpenAPI – where the spec drives the implementation rather than documenting it after the fact – something changes. Documentation generates automatically. SDKs regenerate on spec change. Validation middleware uses the same spec the docs reference. The three artifacts are guaranteed consistent because they come from the same source. As a side effect, time spent on API maintenance drops significantly – but the real win is confidence that your docs, clients, and validation all agree on what the API actually does.

WARNING

The biggest OpenAPI mistake: generating a spec from existing code and calling it “documentation.” That’s archaeology, not specification. The spec should drive the code, not describe it after the fact.

OpenAPI Fundamentals

If you’ve worked with OpenAPI specs before, you know they can get unwieldy fast. A moderately complex API generates a YAML file that scrolls forever, and finding anything requires a mental map of where things live. Understanding the structure matters because it determines whether your spec stays maintainable or becomes a 3,000-line file nobody wants to touch.

The Building Blocks

An OpenAPI spec has four top-level sections that matter:



info for metadata

servers for environment URLs

paths for your actual endpoints

components for reusable definitions

The first two are straightforward – title, version, contact info, and a list of base URLs for production, staging, and local development.

The `paths` section is where you define what your API does. Each path (like `/orders` or `/orders/{orderId}`) contains operations – GET, POST, PUT, DELETE – with their parameters, request bodies, and responses. Here's where things get interesting: instead of defining the same `Order` schema inline on every endpoint that returns one, you define it once in `components/schemas` and reference it with `$ref: "#/components/schemas/Order"`.

This applies to everything. Parameters that appear on multiple endpoints? Define them in `components/parameters`. Error responses that every endpoint might return? `components/responses`. Security schemes? `components/securitySchemes`. One change propagates everywhere.

components-example.yaml

```
1  # Reusable components keep your spec DRY
2  components:
3    parameters:
4      PageParam:
5        name: page
6        in: query
7        schema:
8          type: integer
9          minimum: 1
10         default: 1
11
12    responses:
13      NotFound:
14        description: "Resource not found"
15        content:
```



```
16         application/json:
17             schema:
18                 $ref: "#/components/schemas/Error"
19
20     paths:
21         /orders:
22             get: # HTTP verb being specified
23                 operationId: listOrders # Becomes client.listOrders()
24                 parameters:
25                     - $ref: "#/components/parameters/PageParam"
26                 responses:
27                     "404":
28                         $ref: "#/components/responses/NotFound"
```

Reusable components and \$ref patterns.

Notice the `operationId` field in that example. It seems like metadata, but it becomes the method name in generated SDKs. Name your operation `listOrders` and you get `client.listOrders()`. Name it `get_api_v1_orders_list` and your SDK (Software Development Kit) consumers will hate you. Spend the thirty seconds to pick good names — future you will appreciate it.

Schema Design That Doesn't Fight You

Here's a mistake I see frequently: using the same schema for create, update, and read operations. It seems efficient at first. You have an `Order` schema, you use it everywhere, done. Then reality hits.

When creating an order, clients don't send `id`, `createdAt`, or `status` —the server assigns those. When updating, they can only change certain fields. When reading, they get everything back. One schema can't cleanly express all three cases.

The workaround is `readOnly` and `writeOnly` annotations scattered throughout the schema, plus nullable fields that aren't really nullable in certain contexts. Generated SDKs end up with properties that sometimes exist and sometimes don't, depending on the operation. It's a mess.

The cleaner approach: separate schemas for each operation.



crud-schemas.yaml

```

1  # Separate schemas for create vs response
2  components:
3    schemas:
4      CreateOrderRequest:
5        type: object
6        required: [customerId, items]
7        properties:
8          customerId:
9            type: string
10           format: uuid
11          items:
12            type: array
13            minItems: 1
14          notes:
15            type: string
16            maxLength: 500
17      # No id, status, timestamps - server assigns these
18
19      Order:
20        type: object
21        required: [id, customerId, items, status, createdAt]
22        properties:
23          id:
24            type: string
25            format: uuid
26          customerId:
27            type: string
28            format: uuid
29          status:
30            type: string
31            enum: [pending, confirmed, shipped, delivered, cancelled]
32          createdAt:
33            type: string
34            format: date-time
35
36      UpdateOrderRequest:
37        type: object
38        properties:
39          notes:
40            type: string

```



```
41         maxLength: 500
42     status:
43         type: string
44     enum: [confirmed, cancelled]
45     # Only fields the client can modify - no id, createdAt, customerId
```

Separate schemas for create, update, and response operations.

Yes, there's duplication. `customerId` appears in both schemas. That's fine – it's explicit about what each operation accepts and returns. Generated code becomes cleaner, validation becomes tighter, and nobody has to reason about which fields apply in which context.

Composition Patterns

When you do have legitimate shared structure, OpenAPI provides composition through `allOf`. If every entity in your system has `id`, `createdAt`, and `updatedAt`, define a `BaseEntity` and compose it:

allof-composition.yaml

```
1  components:
2    schemas:
3      BaseEntity:
4        type: object
5        properties:
6          id:
7            type: string
8            format: uuid
9          createdAt:
10           type: string
11           format: date-time
12
13      Customer:
14        allOf:
15          - $ref: "#/components/schemas/BaseEntity"
16          - type: object
17            required: [email, name]
18            properties:
19              email:
20                type: string
```



```

21         format: email
22     name:
23         type: string

```

Using `allOf` to compose shared base properties into entity schemas.

For polymorphic types – where a field can be one of several shapes – use `oneOf` with a discriminator. Payment methods are the classic example: credit cards have `last4` and `expiryMonth`, bank transfers have `routingNumber` and `accountNumber`, PayPal has `email`. The discriminator tells parsers which schema applies based on a `type` field.

#	Pattern	When to Use
1	Separate CRUD schemas	Create/update/read operations have different field requirements
2	<code>allOf</code> composition	Multiple entities share identical base fields (audit timestamps, IDs)
3	<code>oneOf</code> polymorphism	A field can be one of several distinct types with different structures

OpenAPI schema design patterns and when to use them

The `oneOf` pattern in practice:

oneof-polymorphism.yaml

```

1  components:
2    schemas:
3      PaymentMethod:
4        oneOf:
5          - $ref: "#/components/schemas/CreditCard"
6          - $ref: "#/components/schemas/BankTransfer"
7        discriminator:
8          propertyName: type
9      CreditCard:
10       type: object
11       required: [type, last4, expiryMonth]
12       properties:

```




```
13         type:
14             type: string
15             enum: [credit_card]
16         last4:
17             type: string
18         expiryMonth:
19             type: integer
20     BankTransfer:
21         type: object
22         required: [type, routingNumber]
23         properties:
24             type:
25                 type: string
26                 enum: [bank_transfer]
27             routingNumber:
28                 type: string
```

Using oneOf with a discriminator for polymorphic payment methods.

INFO

A good rule of thumb: if you're adding `readOnly` , `writeOnly` , or complex `nullable` logic to make one schema work everywhere, you probably need separate schemas instead.

Spec-First vs Code-First

The spec-first vs code-first debate comes down to one question: what's your source of truth? In spec-first, the OpenAPI document is the contract – you design the API there, then generate server stubs and client SDKs from it. In code-first, your implementation is the contract – you add annotations to your code, then generate the spec from those annotations.

I've worked with both, and my bias is toward spec-first for any API with multiple consumers. Here's why.



The Case for Spec-First

When the spec is the source of truth, several good things happen. API design becomes a distinct activity that happens **before** you start coding. You can share the spec with frontend teams, partner integrators, or product managers and get feedback before you've written a line of implementation code. Breaking changes become visible in the spec diff, not hidden in a code change that looks innocuous.

The workflow looks like this: design the API in your spec editor, review with stakeholders, generate server stubs (which give you the interface to implement), fill in the business logic, then generate client SDKs and documentation from the same spec. Everything stays in sync because everything comes from the same source.

You'll want a decent editor for spec-first work. Here are the main options:

Editor	Best For
Swagger Editor Next	Browser-based editing with real-time validation; supports OpenAPI 3.1 and AsyncAPI
Stoplight Studio	Visual, GUI-driven design without memorizing YAML syntax
42Crunch OpenAPI Editor (VS Code)	IDE-integrated editing with IntelliSense, visual navigation, and "Try it" testing
IntelliJ OpenAPI Plugin	Teams already in JetBrains IDEs; can generate specs from existing controllers
Postman	Teams already using Postman for testing who want spec editing in the same workspace

OpenAPI spec editors for design-first workflows.

The downsides are real too. Generated code rarely matches your team's style preferences perfectly. Every change requires editing the spec first, which adds friction when you're iterating quickly. And spec-first requires discipline – it's tempting to make a quick fix in code and "update the spec later," which is how drift starts.

Spec-first works best for public APIs, APIs consumed by multiple teams, and anywhere you need a formal contract that outlives any particular implementation.



The Case for Code-First

Code-first feels more natural to most developers. You build the API in your language and framework of choice, sprinkle in some annotations or decorators, and out comes a spec. The code is always accurate because the code **is** the specification.

Frameworks like FastAPI (Python), tsoa (TypeScript), and springdoc (Java) make this almost friction-free. You write type hints or interfaces, add decorators for additional metadata, and the spec generates automatically on build.

fastapi-example.py

```
1  # FastAPI: code-first with automatic spec generation
2  from fastapi import FastAPI
3  from pydantic import BaseModel
4  from uuid import UUID
5  from datetime import datetime
6
7  app = FastAPI(title="Order API", version="1.0.0")
8
9  class Order(BaseModel):
10     id: UUID
11     customer_id: UUID
12     status: str
13     created_at: datetime
14
15  @app.get("/orders/{order_id}", response_model=Order)
16  async def get_order(order_id: UUID) -> Order:
17     # Visit /docs for auto-generated Swagger UI, /openapi.json for the spec
18     return await fetch_order(order_id)
```

FastAPI automatically generates OpenAPI specs from Python type hints.

The tradeoffs: spec quality depends entirely on how thorough your annotations are. It's easy to ship an endpoint with minimal documentation because "it works." API design happens **during** coding, which means design decisions get made under time pressure. And breaking changes are easier to introduce accidentally because there's no separate spec review step.

Code-first works well for internal APIs, rapid prototypes, and situations where you're the primary consumer of your own API.



Keeping Spec and Code in Sync

Whichever approach you choose, drift is the enemy. The spec says one thing, the implementation does another, and now your documentation is lying to consumers.

The solution is CI validation. On every pull request, automatically lint the spec, check for breaking changes against the main branch, and run contract tests that verify the implementation matches the spec.

api-validation-steps.sh

```
1  #!/bin/bash
2
3  # The key CI steps for catching drift
4  npx @stoplight/spectral-cli lint openapi.yaml           # Spec quality
5  npx openapi-diff main:openapi.yaml openapi.yaml        # Breaking changes
6  npx dredd openapi.yaml http://localhost:3000           # Contract tests
```

The three validation commands to run in CI. (Full workflow example in the CI/CD section.)

The key tools here: Spectral for linting (catches spec-level issues like missing descriptions or inconsistent naming), openapi-diff for detecting breaking changes (removed endpoints, required fields added, type changes), and Dredd or Schemathesis for contract testing (actually hits your running server and validates responses against the spec).

⚠ WARNING

If you don't validate in CI, spec and implementation **will** drift. It's not a matter of discipline – it's that small, well-intentioned changes accumulate into documentation that actively misleads consumers.

Documentation Generation

Once you have an OpenAPI spec, generating documentation is almost trivially easy. Point a tool at your YAML file, run a command, and out comes a browsable API reference. The hard part isn't generation – it's making documentation that developers actually want to read.



Choosing a Documentation Tool

The three main options cover different use cases:

➤ **Redoc**

Generates static HTML from your spec. You get a three-panel layout (navigation, content, code samples), search, and nested schema expansion. It's a single HTML file you can host anywhere – GitHub Pages, S3, your CDN. No backend required. This is my default choice for most projects because it's zero-maintenance once deployed.

➤ **Swagger UI**

Provides interactive documentation where developers can make live requests against your API. The "Try it out" button is genuinely useful for exploration and debugging. The tradeoff: you need to think about authentication, CORS, and whether you want people hitting your production API from the docs.

➤ **Stoplight Elements**

A React component library for embedding documentation in an existing portal or documentation site. If you're building a developer portal with tutorials, guides, and API reference together, Elements lets you integrate the API docs natively rather than linking out to a separate site.

generate-docs.sh

```
1  #!/bin/bash
2
3  # Redoc: generate static HTML
4  npx @redocly/cli build-docs openapi.yaml -o docs/api-reference.html
5
6  # Swagger UI: serve with Docker
7  docker run -p 8080:8080 \
8    -e SWAGGER_JSON=/spec/openapi.yaml \
9    -v ./openapi.yaml:/spec/openapi.yaml \
10   swaggerapi/swagger-ui
```

Generating docs with Redoc and Swagger UI.

Making Documentation Worth Reading

Generated documentation is only as good as your spec. A minimal spec produces minimal docs – endpoint names, parameter lists, response schemas. Functional, but not helpful.



The difference between documentation that gets ignored and documentation that developers bookmark comes down to three things: descriptions that explain **why** not just **what**, examples that show realistic data, and error responses that help with debugging.

Descriptions support Markdown, so use it. Don't just write "Gets an order"—explain what permissions are required, what the caching behavior is, what the rate limits are. This context belongs in the spec itself, not in a separate wiki that gets out of sync.

rich-descriptions.yaml

```
1  # Descriptions that actually help developers
2  paths:
3    /orders/{orderId}:
4      get:
5        summary: Get order by ID
6        description: |
7          Retrieves a single order by its unique identifier.
8
9          **Authorization**: Requires `orders:read` scope.
10
11          **Rate limiting**: 100 requests/minute per API key.
12
13          **Caching**: Responses cached for 60 seconds.
14          Use `Cache-Control: no-cache` to bypass.
```

OpenAPI descriptions with authorization, rate limits, and caching behavior.

Examples matter even more than descriptions. A schema tells you the **shape** of data; an example shows you what it actually **looks like**. OpenAPI supports multiple named examples per response, so you can show different scenarios — a pending order vs. a shipped order, a successful response vs. an error.

multiple-examples.yaml

```
1  # Multiple examples show different scenarios
2  responses:
3    "200":
4      description: Order retrieved successfully
5      content:
6        application/json:
```



```
7      examples:
8        pending_order:
9          summary: Pending order
10         value:
11           id: "550e8400-e29b-41d4-a716-446655440000"
12           status: pending
13           items: [{productId: "prod-001", quantity: 2}]
14        shipped_order:
15          summary: Shipped order
16         value:
17           id: "660e8400-e29b-41d4-a716-446655440001"
18           status: shipped
19           trackingNumber: "1Z999AA10123456784"
```

Named examples for different order states.

These individual techniques — rich descriptions, multiple examples, error scenarios — compound when applied consistently across your spec. The table below summarizes which documentation elements to prioritize and what each should contain, so you can audit an existing spec or build a new one with the right coverage from the start.

#	Documentation Element	What to Include
1	Operation description	Auth requirements, rate limits, caching behavior, business context
2	Parameter descriptions	Format constraints, valid ranges, default behavior
3	Response examples	Multiple scenarios (success variants, common errors)
4	Schema descriptions	Field meanings, valid values, relationships to other entities

Documentation elements that improve developer experience.



✓ SUCCESS

The time you spend on descriptions and examples pays off in reduced support requests. Developers who can self-serve from good docs don't need to ask you questions.

Good documentation tells developers what your API does. But they still have to write the code to call it – and that's where SDK (Software Development Kit) generation comes in.

SDK Generation

Generating client SDKs from OpenAPI specs sounds like magic until you actually try it. The promise: run a command, get type-safe clients in TypeScript, Python, Java, whatever your consumers use. The reality: generated code that's technically correct but awkward to use, with naming conventions that don't match your codebase and patterns that feel foreign.

The good news is that SDK (Software Development Kit) generation has matured significantly. With the right generator choice and some post-processing, you can produce SDKs that feel handwritten.

Picking a Generator

The SDK (Software Development Kit) generator landscape breaks down into three categories: comprehensive multi-language tools, language-specific tools with better output, and types-only generators.

OpenAPI Generator (the open-source fork of Swagger Codegen)

Supports 50+ languages and is the default choice for teams that need SDKs in multiple languages from the same spec. The TypeScript output is decent, the Python output is usable, and you get consistent behavior across languages. The tradeoff is that "decent" and "usable" aren't "great"—the generated code tends toward verbosity and doesn't always follow language idioms.



Orval

Focuses exclusively on TypeScript but produces notably better output. It generates React Query hooks, Zod validation schemas, and MSW mock handlers out of the box. If your consumers are React developers, Orval's output feels like code they'd actually write.



openapi-typescript

Takes a different approach: it generates only TypeScript types, no runtime code. You get interfaces for your request/response bodies and use your own HTTP client. This is the lightest-weight option and works well when you want type safety without adopting a generated SDK wholesale.



These tools solve different problems, so the right choice depends less on raw feature count and more on who will consume the generated code. A frontend-heavy team may prioritize React Query integration and ergonomic TypeScript output, while a platform team supporting multiple languages may accept less idiomatic clients in exchange for broader coverage. The comparison below makes those tradeoffs explicit.

Generator	Languages	Best For
OpenAPI Generator	50+ (TypeScript, Python, Java, Go, etc.)	Multi-language SDK needs, polyglot organizations
Orval	TypeScript only	React apps, teams wanting React Query/Zod integration
openapi-typescript	TypeScript types only	Type safety without generated runtime code
AutoRest	TypeScript, Python, Java, C#, Go	Azure integrations, Microsoft ecosystem

SDK generators and their sweet spots.

Generating Your First SDK

The basic workflow is straightforward. Install the generator, point it at your spec, choose an output target:



generate-sdk.sh

```
1  #!/bin/bash
2
3  # OpenAPI Generator: TypeScript with Axios
4  npx @openapitools/openapi-generator-cli generate \
5    -i openapi.yaml \
6    -g typescript-axios \
7    -o ./sdk/typescript \
8    --additional-properties=supportsES6=true,npmName=@company/orders-client
9
10 # Orval: TypeScript with React Query
11 npx orval --input openapi.yaml --output ./src/api/generated.ts
```

Basic SDK (Software Development Kit) generation commands.

The `additional-properties` flag is where customization happens. Each generator has dozens of options controlling output structure, naming conventions, and feature flags. The `-g typescript-axios` in the command above tells OpenAPI Generator to use its TypeScript-with-Axios template (one of many available templates – there’s also `typescript-fetch`, `typescript-node`, `python`, `java`, etc.). For the TypeScript-Axios template specifically, the useful properties are `supportsES6` (modern syntax), `withSeparateModelsAndApi` (splits models and API classes into separate files), and `useSingleRequestParameter` (bundles parameters into a single options object).

Making Generated Code Usable

Generated code rarely matches your team’s standards out of the box. Rather than fighting the generator’s templates, embrace post-processing. Generate the code, then transform it.

sdk-pipeline.sh

```
1  #!/bin/bash
2  # Post-process generated SDK
3
4  # Generate raw SDK
5  npx @openapitools/openapi-generator-cli generate \
6    -i openapi.yaml \
7    -g typescript-axios \
8    -o ./sdk/typescript
9
```



```
10 # Format with your project's prettier config
11 npx prettier --write ./sdk/typescript/**/*.ts
12
13 # Run ESLint with auto-fix
14 npx eslint --fix ./sdk/typescript/**/*.ts
15
16 # Remove generator metadata files
17 rm -rf ./sdk/typescript/.openapi-generator
18 rm -f ./sdk/typescript/.openapi-generator-ignore
```

Post-processing pipeline for generated SDKs.

For deeper customization, generators support template overrides. You can export the default Mustache templates, modify them, and pass your custom template directory during generation. This is useful for adding custom headers, changing error handling patterns, or injecting wrapper code. But I'd recommend exhausting post-processing options first – template maintenance becomes a burden when the generator updates.

The pattern I've found works best: generate into a dedicated directory, post-process aggressively, then wrap the generated code with a thin facade that exposes the interface your consumers actually want. The generated code handles serialization, HTTP mechanics, and type definitions. Your wrapper handles configuration, error normalization, and any convenience methods.

INFO

Treat generated SDKs as build artifacts, not source code. Regenerate them in CI, don't commit them to version control, and wrap them with a stable interface your consumers depend on.

Request Validation

Validation is where OpenAPI specs earn their keep. Instead of writing validation logic by hand – checking types, verifying formats, ensuring required fields exist – you can derive it directly from the spec. The same schema that generates your documentation and SDKs also enforces that incoming requests match the contract.



The catch: OpenAPI validation handles **schema** correctness, not **business** correctness. It'll reject a request where `quantity` is a string instead of an integer, or where `email` doesn't match email format. It won't reject a request where `customerId` is a valid UUID that doesn't exist in your database.

Adding Validation Middleware

In Node.js/Express, `express-openapi-validator` is the standard choice. You point it at your spec file, and it automatically validates request bodies, query parameters, path parameters, and headers against your schemas. Invalid requests get rejected before they reach your route handlers.

express-validation.ts

```
1 // Express.js with OpenAPI request validation
2 import express from 'express';
3 import * as OpenApiValidator from 'express-openapi-validator';
4
5 const app = express();
6 app.use(express.json());
7
8 // Add validation middleware - validates against openapi.yaml
9 app.use(
10   OpenApiValidator.middleware({
11     apiSpec: './openapi.yaml',
12     validateRequests: true,
13     validateResponses: true, // Optional: also validate your responses
14   })
15 );
16
17 // Route handlers only see validated requests
18 app.post('/orders', async (req, res) => {
19   // req.body already validated against CreateOrderRequest schema
20   // If validation failed, a 400 was returned before reaching here
21   const order = await createOrder(req.body);
22   res.status(201).json(order);
23 });
```

Express middleware that validates requests against your OpenAPI spec.



The `validateResponses: true` option is worth enabling in development and staging. It catches cases where your implementation returns data that doesn't match your spec – a useful early warning that your spec and code have drifted apart. You might disable it in production for performance reasons, but keeping it on during development catches a lot of bugs.

For Python, FastAPI handles this automatically if you're using Pydantic models. Your type hints **are** your validation schema, and FastAPI generates the OpenAPI spec from them. If you're using a different Python framework, libraries like `openapi-core` provide similar middleware.

Other ecosystems have their own solutions:

Framework	Library	Approach
Rails	<code>committee</code>	Rack middleware that validates against OpenAPI specs
Laravel	<code>spectator</code>	Request/response validation with PHPUnit integration
.NET	<code>Microsoft.AspNetCore.OpenApi</code>	Built-in validation in ASP.NET Core 9+; earlier versions use NSwag
Go	<code>kin-openapi</code>	Middleware for Echo, Chi, Gin; validates requests and responses
Spring Boot	<code>springdoc-openapi</code>	Code-first with automatic spec generation and validation

OpenAPI validation libraries across frameworks.

Handling Validation Errors

Raw validation errors from middleware tend to be developer-friendly but user-hostile. You'll get JSON Pointer paths like `/body/items/0/quantity` and technical messages like `must be >= 1`. These are fine for debugging but not great for API consumers trying to fix their requests.

Transform validation errors into something more useful:



error-handler.ts

```
1 // Transform validation errors into user-friendly format
2 app.use((err: any, req: express.Request, res: express.Response, next:
  express.NextFunction) => {
3   if (err.status === 400 && err.errors) {
4     const formatted = err.errors.map((e: any) => ({
5       field: e.path.replace('/body/', '').replace(/\\/g, '.'),
6       message: humanizeError(e),
7     }));
8
9     return res.status(400).json({
10      code: 'VALIDATION_ERROR',
11      message: 'Request validation failed',
12      errors: formatted,
13    });
14  }
15  next(err);
16 });
17
18 function humanizeError(error: { path: string; message: string; params?: Record<string,
  unknown> }): string {
19   const field = error.path.replace('/body/', '').replace(/\\/(\d+)\\/g,
  '[$1].').replace(/\\/g, '.');
20   // Convert technical messages to readable ones
21   if (error.message.includes('must be >= ')) {
22     return `${field} must be at least ${error.params?.limit ?? 'the minimum'}`;
23   }
24   if (error.message.includes('required')) {
25     return `${field} is required`;
26   }
27   return `${field}: ${error.message}`;
28 }
29
30 // /body/items/0/quantity -> "items[0].quantity must be at least 1"
```

Error handler that transforms technical validation errors into readable messages.

What Validation Catches (and Doesn't)

It's important to understand the boundaries of schema validation:



Validation Catches	Validation Misses
Wrong types (string instead of integer)	Valid types with invalid business meaning
Missing required fields	Fields that are conditionally required
Format violations (invalid UUID, email)	Valid formats that don't exist (nonexistent user ID)
Out-of-range values (quantity < 1)	Domain-specific constraints (can't order discontinued items)
Invalid enum values	Contextually invalid values (can't cancel a shipped order)

Schema validation boundaries.

Schema validation is your first line of defense – it rejects obviously malformed requests before they hit your business logic. But you still need business validation after schema validation passes. A request with `{"customerId": "valid-uuid", "productId": "valid-uuid"}` passes schema validation perfectly, but both IDs might reference entities that don't exist or aren't accessible to the requesting user.

WARNING

Don't confuse schema validation with authorization or business rules. A valid request body doesn't mean an authorized request. Always layer business validation and permission checks on top of schema validation.

CI/CD Integration

Everything we've discussed – documentation generation, SDK (Software Development Kit) generation, validation – becomes reliable only when it's automated. If generating docs requires someone to remember to run a command, docs will fall out of date. If SDK (Software Development Kit) generation is a manual step, it'll be skipped when someone's in a hurry. CI/CD is what turns OpenAPI from “nice to have” into “always accurate.”



The Core Pipeline

A solid OpenAPI pipeline has four stages: lint the spec, check for breaking changes, run contract tests, and generate artifacts. Here's what that looks like in GitHub Actions:

.github/workflows/api.yml

```
1  name: API Pipeline
2  on: [push, pull_request]
3
4  jobs:
5    lint:
6      runs-on: ubuntu-latest
7      steps:
8        - uses: actions/checkout@v4
9        - name: Lint OpenAPI spec
10         run: npx @stoplight/spectral-cli lint openapi.yaml
11
12    breaking-changes:
13      runs-on: ubuntu-latest
14      if: github.event_name == 'pull_request'
15      steps:
16        - uses: actions/checkout@v4
17        with:
18          fetch-depth: 0
19        - name: Check for breaking changes
20          run: |
21            git show origin/${{ github.base_ref }}:openapi.yaml > base.yaml
22            npx openapi-diff base.yaml openapi.yaml --fail-on-incompatible
23
24    contract-test:
25      runs-on: ubuntu-latest
26      needs: [lint]
27      steps:
28        - uses: actions/checkout@v4
29        - run: npm ci && npm start &
30        - run: sleep 10 # Wait for server to be ready
31        - run: npx dredd openapi.yaml http://localhost:3000
32
33    generate:
34      runs-on: ubuntu-latest
```




```
35     needs: [lint, contract-test]
36     if: github.ref == 'refs/heads/main'
37     steps:
38       - uses: actions/checkout@v4
39       - name: Generate SDK
40         run: npx @openapitools/openapi-generator-cli generate -i openapi.yaml -g
              typescript-axios -o ./sdk
41       - name: Generate docs
42         run: npx @redocly/cli build-docs openapi.yaml -o ./docs/api.html
```

GitHub Actions pipeline with linting, breaking change detection, contract tests, and artifact generation.

The key insight: the `breaking-changes` job only runs on pull requests and compares against the base branch. This gives reviewers visibility into API changes before they merge. The `generate` job only runs on main after tests pass, ensuring published artifacts always reflect tested code.

Linting with Spectral

Spectral is the standard linter for OpenAPI specs. It ships with sensible defaults (valid spec structure, required fields, consistent naming) and lets you add custom rules for your organization's standards.

The built-in `spectral:oas` ruleset catches structural issues – invalid syntax, missing required fields, unreferenced schemas. For most teams, that's enough to start:

`.spectral-basic.yaml`

```
1  # Start with the built-in ruleset
2  extends: ["spectral:oas"]
```

Minimal Spectral config using only the built-in OpenAPI rules.

Once you have that in CI, add custom rules to enforce your organization's conventions – naming patterns, required documentation, security definitions:



`.spectral-extended.yaml`

```
1  # Extended config with custom rules
2  extends: ["spectral:oas"]
3
4  rules:
5    # Require descriptions on all operations
6    operation-description:
7      severity: warn
8      given: "$.paths.*[get,post,put,delete]"
9      then:
10       field: description
11       function: truthy
12
13    # Enforce kebab-case paths
14    path-keys-kebab-case:
15      severity: error
16      given: "$.paths"
17      then:
18       function: pattern
19       functionOptions:
20         match: "^(/[a-z][a-z0-9-]*)+$"
```

Extended Spectral config with custom rules for consistency.

Detecting Breaking Changes

Breaking changes are the silent killer of API reliability. Someone adds a required field to a request body, thinking it's a small change. Three downstream services start failing in production because they don't send that field.

`openapi-diff` compares two spec versions and flags incompatible changes:

- Removed endpoints
- Removed or renamed fields
- New required fields in request bodies
- Changed field types
- Narrowed enum values



Run it in PR checks with `--fail-on-incompatible` to block merges that would break consumers. When breaking changes are intentional (major version bumps), you can override the check – but at least it's a conscious decision, not an accident.

Contract Testing

Contract tests verify your implementation actually matches your spec. Dredd and Schemathesis are the main tools:

- **Dredd**
Reads your spec, generates requests for each endpoint, hits your running server, and validates responses against the schema. It's deterministic – same spec produces same tests.
- **Schemathesis**
Takes a property-based testing approach. It generates random valid inputs based on your schemas and looks for edge cases that break your implementation. It's better at finding bugs but less predictable.

I typically use Dredd in CI for consistent, fast feedback and Schemathesis periodically for deeper testing.

✓ SUCCESS

The combination of Spectral (spec quality), openapi-diff (breaking change detection), and Dredd (contract verification) catches most API issues before they reach production. If you only automate one thing, automate this.

Conclusion

The difference between OpenAPI that delivers value and OpenAPI that becomes shelfware comes down to one thing: is the spec the source of truth, or is it an artifact you generate and forget?

When the spec drives everything – documentation, SDKs, validation, contract tests – you get consistency for free. Change the spec, regenerate artifacts, deploy. Documentation can't drift because it's generated. SDKs can't disagree with the server because they come from the same source. Validation can't miss edge cases the spec covers because it reads the spec directly.



The setup cost is real. You'll spend time choosing generators, configuring linters, wiring up CI pipelines. But that cost is paid once. The alternative — manually maintaining documentation, hand-coding SDKs, duplicating validation logic — is paid continuously, and it compounds as your API grows and your consumer base expands.

If you take one thing from this article: treat your OpenAPI spec as infrastructure, not documentation. The spec isn't describing your API — it **is** your API contract, and everything else flows from that.

Copyright © 2023 Webstack Builders, Inc.

The text, diagrams, and images in this work are licensed under CC BY-NC 4.0

All code samples in this article are licensed under the MIT License. Feel free to use, modify, and distribute them in any project.

<https://www.webstackbuilders.com/articles/openapi-spec-documentation-sdk-generation-validation>

